BOUNDED AI FOR DIGITAL-PAYMENT SAFETY

# TrustLens APAC

**Make the safe next step obvious** before a scam becomes a transaction.

THEME // BFSI: INTELLIGENT RISK, FRAUD & FINANCIAL EXPERIENCES

TL

# The moment before the loss

A scam is not just a bad message. It is a high-pressure decision.

> “Your refund is pending. Share your OTP and click the link to release the money.”

One message. Four signals. Zero context for the user.

| | |
|---|---|
| 01<br>**Impersonation**<br>Looks like a trusted institution. | 02<br>**Urgency**<br>Compresses time for verification. |
| 03<br>**Credential ask**<br>Requests an OTP or password. | 04<br>**Redirect**<br>Moves the user to an unsolicited link. |

THE GAP // Users need an immediate safe action. Analysts need evidence, policy context, and a review boundary.

# From confusing message to safe action

TrustLens turns unstructured evidence into a conservative, explainable intervention.

01 / INTAKE

## Capture

Paste a message or upload a screenshot. Select the user's language.

→

02 / EVIDENCE

## Normalize

Gemini extracts sender claims, URLs, urgency, requested actions, and evidence IDs.

→

03 / REASON

## Ground

Pattern analysis and policy retrieval produce a schema-constrained risk synthesis.

→

04 / INTERVENE

## Protect

Safety guard checks the recommendation. High-risk or uncertain cases route to human review.

# A decision-support layer, not another chatbot

Use generative reasoning where it creates leverage. Use deterministic software where safety matters.

01 / EVIDENCE-FIRST

## Every claim has a trail

Risk rationales point to extracted evidence IDs, not an unexplained score.

02 / POLICY-GROUNDED

## Advice has a source

Recommendations are tied to versioned safety policies and prohibited actions.

03 / SAFETY-INVARIANT

## Model output is not the last word

A Python guard can override unsafe guidance involving OTPs, links, and remote access.

04 / HUMAN-BOUND

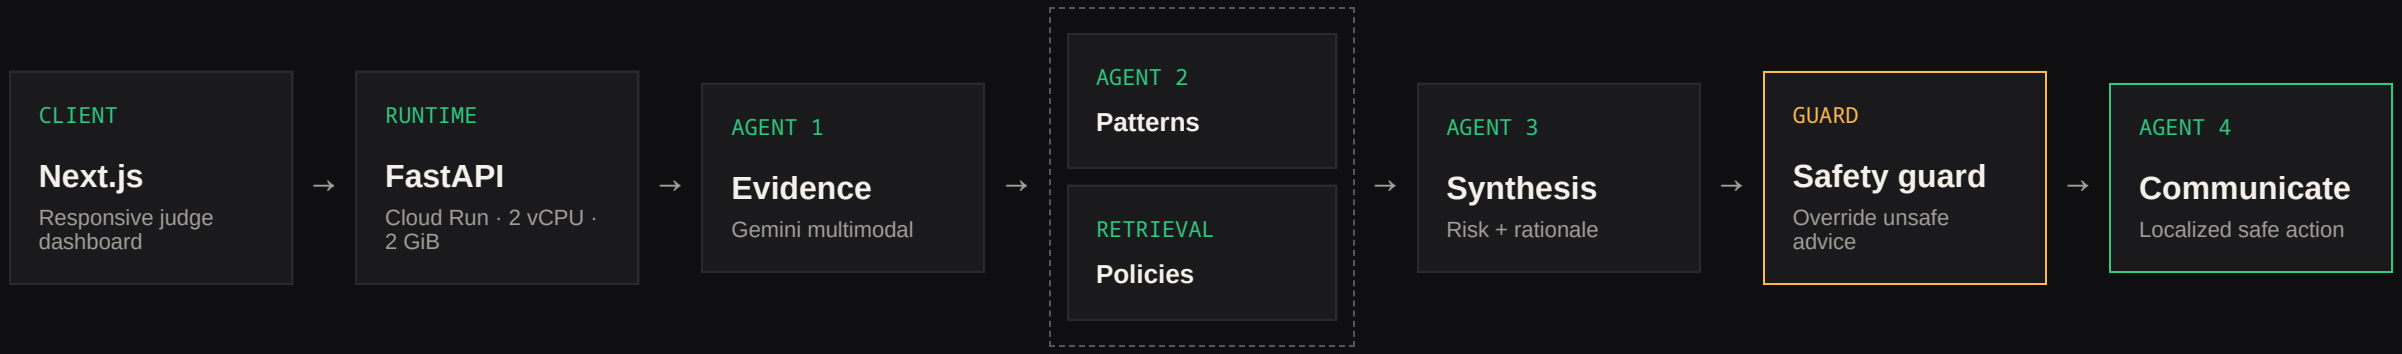
## Know when to stop

High-risk, uncertain, or unsupported cases are explicitly routed for review.

THE THESIS // TRUST IS WON BY SHOWING THE USER WHAT TO DO — AND WHAT NOT TO DO.

# Bounded agents. Google Cloud backbone.

A server-side pipeline with explicit contracts, parallel reasoning, and a deterministic safety boundary.



# The product judges can use

One input. One explainable result. One safer next step.

INPUT

## Message or screenshot

Text and multimodal evidence enter the same workflow.

OUTPUT

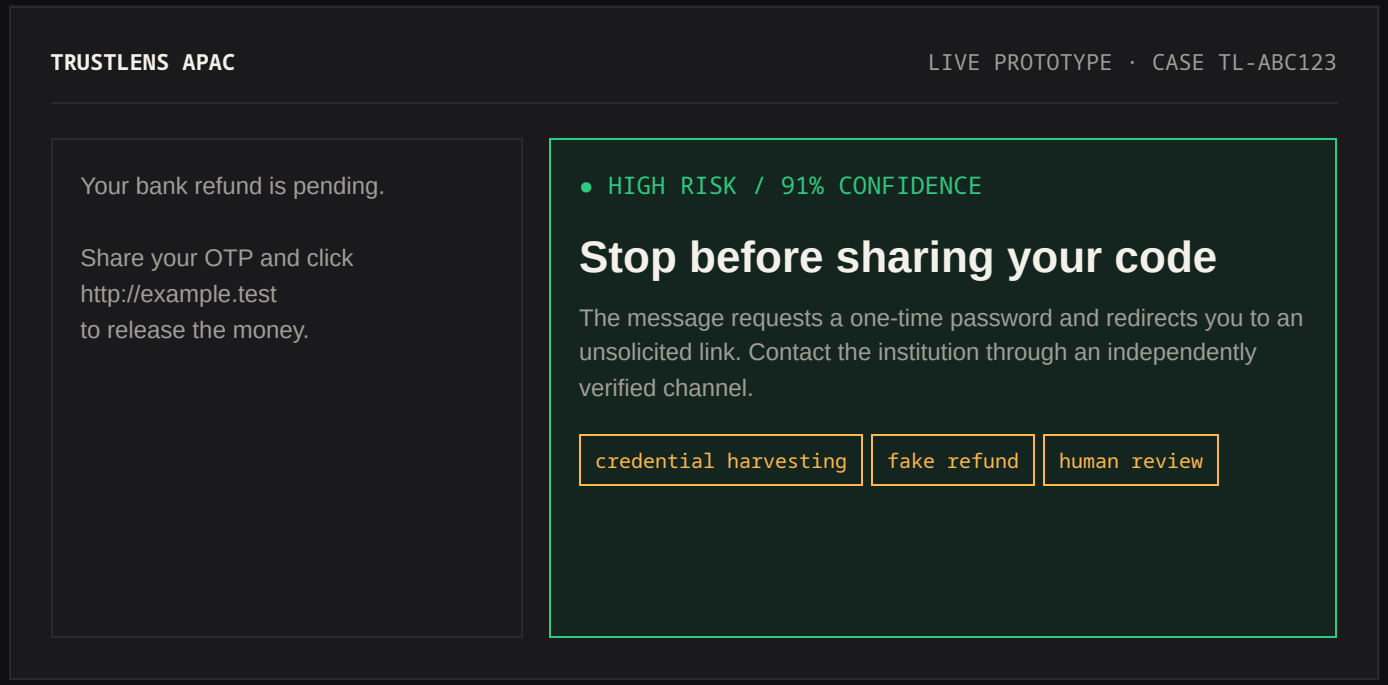
## Risk + rationale

Risk band, confidence, patterns, evidence IDs, and uncertainty.

HANDOFF

## Human review

Analyst can correct or approve the next step; no autonomous account action.



# Where Gemini creates real leverage

Generative reasoning is visible, bounded, and testable at every stage.

MULTIMODAL

## Read what users actually receive

Gemini extracts text and visual cues from screenshots, including sender claims, links, and urgency signals.

SCHEMA

## Return objects, not prose

Pydantic contracts constrain every agent response and make the UI and guard path deterministic.

PARALLEL

## Separate reasoning responsibilities

Pattern matching and policy analysis run independently after extraction, reducing coupling and latency.

LOCALIZE

## Make safety actionable

Communication is generated in the selected language, while prohibited-action warnings remain explicit.

# Designed to measure what matters

Prototype targets are explicit. Observed results are reported only after a real synthetic run.

| ≥90% | 0% | 100% |
| --- | --- | --- |
| **High-risk recall target** | **Unsafe-action target** | **Review routing target** |
| High cases identified as high or review | Prohibited advice in evaluated outputs | Ambiguous/high-risk cases escalated |

## Synthetic evaluation set

30+ versioned cases across six scam patterns, text and screenshot inputs, multiple languages, ambiguous cases, and adversarial safety tests.

NO INVENTED PERFORMANCE CLAIMS // TARGETS ARE NOT PRODUCTION GUARANTEES

## Expansion path

01 · wallets

02 · banks

03 · telecom fraud

04 · remittances

# Safe by design

TrustLens assists interpretation. It does not perform irreversible financial actions.

01 / ALWAYS SHOW

## Evidence + uncertainty

- Risk rationale linked to evidence IDs.
- Grounding sources exposed in audit metadata.
- Uncertainty remains visible to the user.

02 / ALWAYS ESCALATE

## Human review

- High-risk cases require review.
- Missing evidence forces review.
- Analyst correction and approval are explicit.

03 / NEVER DO

## Autonomous action

- Never move money or freeze an account.
- Never request an OTP, PIN, or password.
- Never contact a bank without an authorized workflow.

THE SAFE NEXT STEP IS A PRODUCT FEATURE

# Make the safe next step obvious.

**TrustLens** brings grounded Gemini reasoning, deterministic safety, and human oversight to the moment before a scam becomes a transaction.

LIVE PROTOTYPE / CLOUD RUN

GITHUB + README

3-MINUTE DEMO